

# High Performance Computing – Course 4: OpenMP – Task-based parallel runtime for shared memory machines

Jonathan Rouzaud-Cornabas

LIRIS / Insa de Lyon – Inria Beagle

# Références

- Cours OpenMP, F. Roch (Grenoble)
- Cours OpenMP, J. Chergui & P.-F. Lavalée (IDRIS)
- <http://www.openmp.org>
- <http://www.openmp.org/mp-documents/spec30.pdf>
- <http://www.idris.fr>
- <http://ci-tutor.ncsa.illinois.edu/login.php>
- Using OpenMP , Portable Shared Memory Model, Barbara Chapman


# A bit of History

- Multitasking has been around for a long time, but specific libraries / languages
- Increase of multi-core computers (shared memory machines) → requirement for a standard
- OpenMP-1 (1997), OpenMP-2 (2000), OpenMP-3 (2008), OpenMP-4.0 (2013), OpenMP-4.5 (2015), OpenMP-5.0 (2018), OpenMP-5.1 (2020+)
- GPU support in OpenMP ( $\approx 10\%$  of loss)

# Multi-tasking programming model on shared memory architecture

- Several tasks are executed in parallel
- Memory is shared (physically or virtually)
- Communication between tasks is done by reads and writes in the shared memory.
- Example: general multi-core processors share a common memory, tasks can be assigned to separate “cores”


# Multitasking programming on UMA architectures



Memory is shared

- Uniform memory access (UMA) architectures
- An inherent problem: memory constraints


# Multitasking programming on NUMA architectures



Memory is directly attached to multi-core chips

- Non-uniform memory access (NUMA) architectures


# Multitasking programming on NUMA architectures



**ACCES DISTANT**  
**ACCES LOCAL**

# General Concepts

- A OpenMP program = A process with threads
- Consequence: Variable shared or not
- Sequential (always on thread 0) and parallel (multiple threads) composition







# Type of Parallelization

Work sharing essentially consists of:


- execute a loop by distributing the iterations between the tasks
- execute multiple sections of code but only one per task
- execute multiple occurrences of the same procedure by different tasks (orphaning)



Boucle parallèle




Sections parallèles



Procédure parallèle (orphaning)

# Synchronization


- It is sometimes necessary to introduce a synchronization between the concurrent tasks to avoid, for example, that these modify in any order the value of the same shared variable (case of reduction operations).



# Distribution of Tasks

Tasks are assigned to processors by the operating system. Different cases can occur:

- at best, at each moment, there is one task per processor with as many tasks as dedicated processors during the whole execution time
- at worst, all tasks are processed sequentially by one and only one processor
- in reality, for essentially reasons of operation on a machine whose processors are not dedicated, the situation is generally intermediate.



# Features of the OpenMP model

- **Benefits**

- Transparent and portable “ thread ” management
- Easy programming

- **Drawbacks**

- Data locality problem
- Shared but not hierarchical memory
- Efficiency not guaranteed (impact of the material organization of the machine)
- Limited scaling, moderate parallelism

# OpenMP VS Threads


Two methods of reflecting the distribution

- A task: a set of calculations (common concept)
  - **Thread**: We explicitly define the distribution of tasks between threads
  - **Thread**: More complex code (we must write thread management)
  - **OpenMP**: We don't define the distribution
  - **Thread**: We can choose the order of execution
  - **OpenMP**: Non-deterministic execution order
- 
- **Thread**: Static scheduling / at the expense of the developer
  - **OpenMP**: Dynamic scheduling / at the expense of the system

# OpenMP VS MPI


These are two complementary models of parallelization.

- OpenMP and MPI has a C and C++ interface
- MPI is a multiprocess model whose communication mode between the processes is explicit (the management of communications is the responsibility of the user)
- OpenMP is a multitasking model whose communication mode between tasks is implicit (communication management is the responsibility of the compiler).



## How it works ?

- It is the developer's responsibility to introduce OpenMP directives into his code (at least in the absence of automatic parallelization tools)
- When the program is executed, the operating system builds a parallel region on the “fork and join” model
- When entering a parallel region, the master task creates / activates (fork) “child” processes (light processes) which disappear / doze off at the end of the parallel region (join) while the task master continues the execution of the program alone until the entry of the next parallel region



# Directives/pragmas format

- Sentinelle directive [clause[clause]..]

```
#pragma omp parallel private(a,b) firstprivate  
{  
    ...  
}
```

- The line is interpreted if the OpenMP flag is activated when calling the compiler otherwise it is a comment → portability



# Construction of a parallel region

```
#include <omp.h>

int main(int argc, char** argv) {
    int a,b,c;

    /* Sequential code executed by the master */

    #pragma omp parallel private(a,b) \
        shared(c)
    {
        /* Parallel area executed by all
           threads */
    }

    /* Sequential code */
}
```

# IF clause of the PARALLEL directive

- Conditional creation of a parallel clause region **IF(logic\_expression)**

```
#include <omp.h>
```

```
int main(int argc, char** argv) {  
    int a,b,c;
```

```
    /* Sequential code executed by the master */
```

```
    #pragma omp parallel for if(para_low)
```

```
    {
```

```
        /* Parallel area executed by all  
        threads */
```

```
    }
```

```
    /* Sequential code */
```

```
}
```

- L'expression logique sera évaluée avant le début de la région parallèle.

# OpenMP Threads

- **Defining the number of threads**

- Through an environment variable `OMP_NUM_THREADS`
- Through a function : `OMP_SET_NUM_THREADS()`
- Through a clause `NUM_THREADS()` of the `PARALLEL` directive

- **Threads are numbered**

- The number of threads does not necessarily equal the number of physical cores
- The thread of index 0 is the master task
- `OMP_GET_NUM_THREADS()` : number of threads
- `OMP_GET_THREAD_NUM()` : index of a thread
- `OMP_GET_MAX_THREADS()` : maximum number of threads

# Include, Compilation and Execution

- `#include <omp.h>`
- `gcc/g++ -fopenmp -o prog prog.c[pp]`
- `export OMP_NUM_THREADS=2 ; ./prog`

# Status of a variable

- The status of a variable in a parallel area is either
  - be SHARED, it is in the global memory
  - is PRIVATE, it is in the stack of each thread, its value is undefined at the entry of the zone
  - or FIRSTPRIVATE, it is in the stack of each thread, its value is defined at the entry of the zone (initial value of the variable)
  - or LASTPRIVATE, it is in the stack of each thread, its value is kept at the output (the value of the last task)
- Declare the status of a variable
  - `#pragma omp parallel PRIVATE(list)`
  - `#pragma omp parallel FIRSTPRIVATE(list)`
  - `#pragma omp parallel SHARED(list)`
- Declare default status `#pragma omp default(SHARED|NONE)`

# PARALLEL directive clauses

- NONE : Any variable must have an explicitly defined status
- SHARED (`variables_list`) : Variables shared between threads
- PRIVATE (`variables_list`) : Private variables for each of the threads, undefined outside the PARALLEL block
- FIRSTPRIVATE (`variables_list`) : Variable initialized with the value that the original variable had just before the parallel section

- **THREADPRIVATE :**
  - A global variable, a file descriptor or static variables (in C)
  - The variable instance persists from one region parallel to the other (unless dynamic mode is active)
  - The instance of the variable in the sequential area is also that of thread 0
- **COPYIN :** allows to transmit the value of the shared variable to all tasks

# Memory allocation

- The default compiler option is generally PRIVATE: local variables allocated in the private  $\Rightarrow$  stack, but some options allow you to change this default and it is recommended not to use these options for OpenMP
- A memory allocation or deallocation operation on a private variable will be local to each task
- If a memory allocation / deallocation operation relates to a shared variable, the operation must be performed by a single task.



# Memory allocation

- The size of the stack is limited, different environment variables or functions can act on this size
- The stack has a size limit for the shell (variable depending on the machine). (`ulimit --s`) (`ulimit --s unlimited`), values expressed in KB.
- OpenMP: OMP environment variable `_STACKSIZE`: defines the number of bytes that each OpenMP thread can use for its private stack

# Work sharing

- Distribution of a loop between threads (loop `//`)
- Distribution of several sections of code between threads, one section of code per thread (sections `//`)
- Execution of a portion of code by a single thread
- Execution of several occurrences of the same procedure by different threads

# Scope of a parallel region

- The reach of a parallel region extends
  - to the code contained lexically in this region (static scope)
  - to the code of subroutines called
- The union of the two represents the dynamic range

# Work Sharing

Guidelines for controlling the distribution of work, data and the synchronization of tasks within a parallel region :

- FOR
- SECTIONS
- SINGLE
- MASTER
- WORKSHARE (only in Fortran)




## Work sharing : Parallel loop

FOR directive: parallelism by distribution of iterations of a loop.

- The mode of distribution of the iterations can be specified in the SCHEDULE clause (coded in the program or thanks to an environment variable)
- A global synchronization is performed at the end of construction END FOR (except if NOWAIT)
- Possibility of having several FOR constructions in a parallel region.
- Loop indices are whole and private
- Infinite and do while loops cannot be parallelized


# Work sharing : SCHEDULE clause

- `#pragma OMP FOR SCHEDULE(STATIC, packet_size)`
  - By default  $packet\_size = \frac{nb\_iterations}{nb\_threads}$
  - Example : 16 iterations (0 à 15), 4 threads : default packet size is 4
- `#pragma OMP FOR SCHEDULE(DYNAMIC, packet_size)`
  - Packets are dynamically distributed to free threads
  - All packages have the same size except possibly the last one, by default the size of the packages is 1.
- `#pragma OMP FOR SCHEDULE(GUIDED, packet_size)`
  - $packet\_size$  : minimum packet size (1 by default) except the last.
  - Maximum packet size at the start of the loop (here 2) then decreases to balance the load.




# Worksharing : SCHEDULE clause


Ex: 24 itérations, 3 threads



Mode **static**, avec  
Taille paquets=nb itérations/nb threads



Glouton : **DYNAMIC**




Cyclique : **STATIC**



Glouton : **GUIDED**

# Aevol and OpenMP

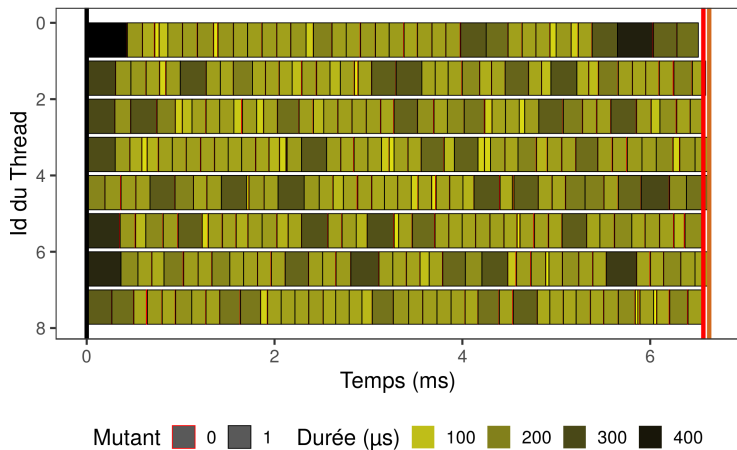
- 3 parallel loop pragma, 2 atomics (counter)




- We also parallelize statistics and checkpointing using task with dependencies
- We try to use task with dependencies everywhere, performance were very bad



# Worksharing : SCHEDULE




# Worksharing : SCHEDULE




# Limitations of the Aeol OpenMP implementation

- The speedup is pretty bad
- Aeol is memory bound (looking for motifs)




- But there is another issue


# Task runtime distribution at a given timestep




# Task runtime during the evolution



# Task runtime with different simulation parameters




# Irregularity and OpenMP



Mutant 0 1 Durée ( $\mu$ s) 100 200 300 400


# Irregularity and OpenMP



8x more cores, 4.4x speedup



# OpenMP scheduling methods



Ordonnancement	Speed Up	Inactivité
Static	33.2	48.1 %
Dynamic	51.2	20 %
Borne Inf	64	0 %

## A “simple” list-scheduling issue

- List scheduling (OpenMP dynamic) :  $(2 - 1/m)$ -approximation,  $O(n \log m)$
- LPT (Longest Processing Time first) :  $(4/3 - 1/(3m))$ -approximation,  $O(n \log n + n \log m)$

## A “simple” list-scheduling issue

- List scheduling (OpenMP dynamic) :  $(2 - 1/m)$ -approximation,  $O(n \log m)$
- LPT (Longest Processing Time first) :  $(4/3 - 1/(3m))$ -approximation,  $O(n \log n + n \log m)$
- But can we model task runtime ?

## A “simple” list-scheduling issue

- List scheduling (OpenMP dynamic) :  $(2 - 1/m)$ -approximation,  $O(n \log m)$
- LPT (Longest Processing Time first) :  $(4/3 - 1/(3m))$ -approximation,  $O(n \log n + n \log m)$
- But can we model task runtime ?
- We try but too difficult and some hardware artifacts


## A “simple” list-scheduling issue

- List scheduling (OpenMP dynamic) :  $(2 - 1/m)$ -approximation,  $O(n \log m)$
- LPT (Longest Processing Time first) :  $(4/3 - 1/(3m))$ -approximation,  $O(n \log n + n \log m)$
- But can we model task runtime ?
- We try but too difficult and some hardware artifacts


## A “simple” list-scheduling issue

- List scheduling (OpenMP dynamic) :  $(2 - 1/m)$ -approximation,  $O(n \log m)$
- LPT (Longest Processing Time first) :  $(4/3 - 1/(3m))$ -approximation,  $O(n \log n + n \log m)$
- But can we model task runtime ?
- We try but too difficult and some hardware artifacts
- We do not need to predict runtime only to predict tasks order

# How to predict task order with Aevol



# Preliminary results



Ordonnancement	Speed Up	Inactivité
Dynamic	51.2	20 %
LDNA	59.8	3.4 %*
LPT	62.8	1.9 %



## Worksharing : SCHEDULE clause

- The choice of the distribution method can be postponed during the execution of the code with `SCHEDULE(RUNTIME)`
- Taking into account the environment variable `OMP_SCHEDULE`
- Example : `export OMP_SCHEDULE='DYNAMIC,400'`

# Reduction : why ?

In sequential

```
for (int i = 0;
     i < N; i++) {
    X=X+a(i)
}
```

In parallel

```
#pragma omp for shared(X)
for (int i = 0;
     i < N; i++) {
    X=X+a(i)
}
```

Reduction : associative operation applied to shared scalar variables

- Each task calculates a partial result independently of the others. Intermediate reductions on each thread are visible locally.
- Then the tasks are synchronized to update the final result in a global variable, by applying the same operator to the partial results.
- Attention, no guarantee of identical results from one execution to another, the intermediate values can be combined in a random order

# Reduction in practice

- Example: `#pragma omp for reduction (op: list)` (op is an operator or an intrinsic function)
- The variables in the list must be shared in the area surrounding the directive!
- A local copy of each variable in the list is assigned to each thread and initialized according to the operation (e.g. 0 for +, 1 for \*)
- The clause will apply to the variables in the list if the instructions are one of the following types:
  - `x = x operator expr`
  - `x = expr operator x`
  - `x = intrinsic (x, expr)`
  - `x = intrinsic (expr,x)`
- x is a scalar variable
- expr is a scalar expression not referencing x
- intrinsic = MAX, MIN, IAND, IOR, IEOR
- operator = +, \*, .AND., .OR., .EQV., .NEQV.

# Reduction in practice

```
int factorial(int number)
{
    int fac = 1;
    #pragma omp parallel for reduction(*:fac)
    for(int n=2; n<=number; ++n)
        fac *= n;
    return fac;
}
```

# Ordered execution : ORDERED

- Run a zone sequentially
  - For debugging purpose
  - For having ordered I/O
- Clause and Directive :  
ORDERED
- The order of execution of instructions in the area framed by the directive will be identical to that of sequential execution, i.e. in the order of iterations

```
#pragma omp for ordered \\  
                        schedule(dynamic)  
for(int n=0; n<100; ++n)  
{  
    files[n].compress();  
  
    #pragma omp ordered  
    send(files[n]);  
}
```

# Unfolding nested loops : COLLAPSE directive


- The COLLAPSE(N) clause allows you to specify a number of loops to unfold to create a large space of iterations
- Loops must be perfectly nested
- Example: If the loops in i and j can be parallelized, and if N and M are small, we can thus parallelize on the whole work corresponding to the 2 loops

```
#pragma omp parallel for\\  
                        collapse(2)  
for(int i=0; i<N; ++i)  
  for(int j=0; j<M; ++j)  
  {  
    tick(i,j);  
  }
```

# Work sharing : parallel SECTIONS parallèles

- **Goal** : Distribute the execution of several independent pieces of code on different tasks
- **A section**: a portion of code executed by one and only one task
- Directive SECTION within a construction SECTIONS

```
#pragma omp sections  
{  
    #pragma omp section  
    { sub0() }  
    #pragma omp section  
    { sub1() }  
    #pragma omp section  
    { sub2() }  
}
```



# Work sharing : exclusive execution

## Construction SINGLE

- Execution of a portion of code by one and only one task (in general, the first which arrives on construction)
- clause NOWAIT : allows not to block the other tasks which by default await its completion
- accepted clauses : PRIVATE, FIRSTPRIVATE
- COPYPRIVATE(*var*): update private copies of the variable (*var*) on all tasks

```
#pragma omp parallel
{
    Work1();
    #pragma omp single
    {
        Work2();
    }
    Work3();
}
```



# Work sharing : exclusive execution

## Structure SINGLE

- Execution of a portion of code by the master task alone
- No synchronization, neither at the beginning nor at the end (unlike SINGLE)
- Beware of variable updates that would be used by other threads
- No clause

```
#pragma omp parallel
{
    Work1();

    #pragma omp master
    {
        Work2();
    }

    Work3();
}
```

# OpenMP and SIMD

Since OpenMP 4.x, it is possible to give hint to compiler through OpenMP pragma to improve automatic vectorization

```
#pragma omp simd reduction(+:sum) aligned(a:64)
for (i = 0; i < num; i++) {
    a[i] = b[i] * c[i];
    sum = sum + a[i];
}
```

# OpenMP and SIMD: WARNING !!!

- Using OpenMP SIMD pragma will bypass the compiler analysis !!!  
(On Intel Compiler, it can refuse/desactivate OpenMP SIMD if it predicts too much performance prediction)
- Use with caution !
- Incorrect results possible !
- Poor performance possible !
- Memory errors possible !

# OpenMP and SIMD: Aligned Memory

- C

```
_mm_malloc(8*sizeof(float),64);
```

```
posix_memalign(pointer,size,align);
```

- C++

- Write a custom allocator
- You need to overload new and delete method within your class

# OpenMP and SIMD

You can mix parallel threads (classic OpenMP) with OpenMP SIMD to use both (you should whenever it is possible)

```
#pragma omp parallel for simd
for (i = 0; i < num; i++) {
    sum = sum + a[i];
}
```

# OpenMP, SIMD and functions

You can declare that a function can be used within a vectorized loop

```
#pragma omp declare simd
float myfunction(float a, float b, float c) {
    return a*b+c;
}

int main(int argc, char** argv) {
    ...
#pragma omp simd
    for(i=0; i < num; i++) {
        OUT[i] = myfunction(array_a[i], array_b[i], array_c[i])
    }
    ...
}
```

# Work sharing : the TASK construct

- A TASK in the OpenMP sense is a unit of work whose execution can be deferred (or start immediately)
  - Allow dynamic generation of tasks
- Allows you to parallelize irregular problems
  - unbounded loops
  - recursive algos
  - producer / consumer diagrams
- A TASK is composed
  - of a code to execute
  - of an associated data environment

# Work sharing : the TASK construct

- The TASK construct explicitly defines an OpenMP TASK
- If a thread encounters a TASK construction, a new instance of the TASK is created (code package + associated data)
- The thread can either execute the task or delay its execution. The task can be assigned to any thread on the team.

```

void increment_list_items
    (node* head)
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            for (node*
                p = head; p;
                p = p->next)
            {
                #pragma omp task
                process(p);
            }
        }
    }
}

```

// p: firstprivate by default



# Work sharing : the TASK construct

- **Note:** the concept existed before construction
- A thread that encounters a construction `texttt PARALLEL`
  - creates a set of implicit TASKs (code + data package)
  - creates a team of threads
  - implicit TASKs are tied to threads, one for each team thread
- When is the TASK executed?
  - The execution of a generated TASK can be assigned, by the scheduler, to a team thread that has finished its work.
  - Standard 3.0 imposes rules on the execution of TASK
  - Example: The pending TASK in the region `//`, must all be executed by the team threads that meet
    - A BARRIER (implicit or explicit)
    - A directive `TASKWAIT`

# Synchronization

- Synchronization of all tasks on the same level of education (global barrier)
- Scheduling of concurrent tasks for the consistency of shared variables (mutual exclusion)
- Synchronization of several tasks among a set (lock mechanism)

# Synchronization: global barrier

- By default at the end of parallel constructions, in the absence of NOWAIT
- Directive BARRIER:  
Explicitly imposes a synchronization barrier:  
each task waits for the end of all the others

```
#pragma omp parallel
{
    /* All threads are running. */
    SomeCode();
    #pragma omp barrier
    // All threads are running but
    // they wait for everyone else before
    SomeMoreCode();
}

#pragma omp parallel
{
    #pragma omp for
    for(int n=0; n<10; ++n)
        Work();
    // Implicit barrier, all threads
    // are waiting for the end of FOR
    SomeMoreCode();
}
```


# Synchronization

It may be necessary to introduce a synchronization between concurrent tasks to prevent them from modifying the value of a variable in any order

Ex :

2 threads ont un espace de mémoire partagé

Espace partagé



Thread 1



Thread 2

# Synchronization: critical regions

## Directive CRITICAL

- It applies to a portion of code

- Tasks execute the critical region in a non-deterministic order, one at a time

- Guarantee threads mutual exclusion access

- Its scope is dynamic

```
int main(int argc, char** argv) {  
    int x;  
    x = 0;  
    #pragma omp parallel shared(x)  
    {  
        #pragma omp critical  
        x = x + 1;  
    } /* end of parallel section */  
}
```

# Synchronization : atomic update

The ATOMIC directive only applies when updating a memory location

- A shared variable is read or modified in memory by only one thread at a time

- Acts on the instruction immediately following if it is in the form :

- $x = x \text{ (op) exp}$
- or  $x = \text{exp (op) } x$
- or  $x = f(x, \text{exp})$
- or  $x = f(\text{exp}, x)$
- op : +, -, \*, /, .AND., .OR., .EQV., .NEQV.
- f : MAX, MIN, IAND, IOR, IEOR
- x is a scalar variable

```
#pragma omp atomic
count = count+1;
```

# FLUSH Directive

- Shared variable values may temporarily stay in registers for performance reasons
- The FLUSH directive guarantees that each thread has access to the values of the shared variables modified by the other threads.

```
/* presumption: int a = 0

/* First thread */
/* Second thread */
    b = 1;
a = 1;
    #pragma omp flush(a,b)
#pragma omp flush(a,b)
    if (a == 0)
if (b == 0)
    {
        /* Critical section */
/* Critical section */
    }
}
```

# FLUSH Directive

- FLUSH implicit in certain regions
  - At the level of a BARRIER
  - At the entry and exit of a region PARALLEL, CRITICAL, or ORDERED, and of a region PARALLEL of work sharing
  - When calling the " lock " functions
  - Immediately before or after each TASK scheduling point
  - When entering and leaving regions ATOMIC (applies to variables updated by ATOMIC)
- No implicit FLUSH
  - At the entrance to a work sharing region
  - When entering or leaving a region MASTER



# Performance and work sharing

- Minimize the number of parallel regions
- Avoid leaving a parallel region to immediately recreate it

```
#pragma omp parallel
{
    #pragma omp for
    for(int n=0; n<10; ++n)
        printf("□%d", n);
}
printf(".\n");
```


```
#pragma omp parallel for
for(int n=0; n<10; ++n)
    printf("□%d", n);
printf(".\n");
```

# Performance and work sharing

- Introduce a `PARALLEL FOR` in loops capable of executing iterations in `//`
- If there are dependencies between iterations, try to remove them by modifying the algorithm
- If there are dependent iterations, introduce constructions `CRITICAL` around the variables concerned by the dependencies
- If possible, group several structures in a single parallel region `FOR`
- If possible, parallel the outermost loop
- Adapt the number of tasks to the size of the problem to be treated in order to minimize the additional costs of task management by the system
- Use `SCHEDULE (RUNTIME)` if necessary
- `ATOMIC REDUCTION` is more efficient than `CRITICAL`

## Performance : False Sharing effects

The consistency of the caches and the negative effects of “false sharing” can have a strong impact on performance



An operation to load a shared cache line invalidates the other copies of this line.

## Performance : False Sharing effects

- The use of structures in shared memory can induce a decrease in performance and a strong limitation of extensibility
  - For performance reasons, use of the cache
  - If several processors handle different but adjacent data in memory, the update of individual elements can cause a complete loading of a cache line, so that the caches are consistent with the memory
- False sharing degrades performance when all of the following conditions are met
  - Shared data is changed on  $\#$  cores
  - Multiple threads on  $\#$  cores update data in the same cache line
  - These updates occur very frequently and simultaneously

# Performance : False Sharing effects

- When shared data is only read, it does not generate false sharing
- In general, the phenomenon of false sharing can be reduced by
  - By possibly privatizing variables
  - Sometimes by increasing the size of the tables (problem size or artificial increase) or by doing padding
  - Sometimes by modifying the way the iterations of a loop are shared between threads (increase the size of the packets)

# Performance : False Sharing effects

```
int a[nthreads];  
  
#pragma parallel for shared(nthreads,a) schedule(static,1)  
for (int i=0; i < nthreads; i++) a[i] = i
```

- **Nthreads** : number of threads executing the loop
- Suppose each thread has a copy of textbf a in its local cache. The packet size of 1 causes a phenomenon of false sharing with each update
- If a cache line can contain C elements of the vector textbf a, we can solve the problem by artificially extending the dimensions of the array (“array padding”): we declare an array **a (C , n)** and we replace **a (i)** with **a (1, i)**

# Performance on multi-core architectures: efficiently schedule threads

- Maximizing the performance of each CPU is highly dependent on the location of memory accesses
- Variables are stored in a memory (and cache) area when they are initialized
  - parallelize the initialization of elements (e.g., array) in the same way (same mode and same packet size) as the job
- The ideal would be to be able to specify thread placement constraints according to thread / memory affinities

# Performance on multi-core architectures: problematic of the diversity of architectures

- types of processors
- numbers of cores
- cache levels
- memory architecture
- types of interconnection of the different components
- Optimization strategies which may differ from one configuration to another



## Performance: conditional parallelization

- Use the IF clause to set up conditional parallelization
- Example: only parallelize a loop if its size is large enough

# OpenMP

- Requires a multi-processor shared memory machine
- Relatively easy to use, even in a sequential program
- Allows progressive parallelization of a sequential program
- The full potential of parallel performance lies in parallel regions
- Within these parallel regions, work can be shared using loops and parallel sections. But we can also differentiate a task for a particular job
- Explicit global or point-to-point synchronizations are sometimes necessary in parallel regions
- Special care must be taken when defining the status of the variables used in a construction

# OpenMP versus MPI

- OpenMP uses memory common to all processes. All communication is done by reading and writing in this memory (and using synchronization mechanisms)
- MPI is a library of routines allowing communication between different processes (located on different machines or not), communications are made by sending or receiving explicit messages
- For the programmer: rewrite the code with MPI, simple addition of directives in the sequential code with OpenMP
- Possibility of mixing the two approaches in the case of a machine cluster with multi-core nodes
- Choice strongly dependent: on the machine, the code, the time that the developer wants to devote to it and the gain sought
- Superior extensibility with MPI